

# Virtual Memory and Memory Management

David Marchant

Based on slides by Troels Henriksen

2022-10-04

Inspired by slides by Randal E. Bryant and David R. O'Hallaron.

Simple memory system example

The Linux virtual memory system


Memory mapping

Dynamic allocation


# Addressing

- 14-bit virtual addresses.
- 12-bit physical addresses.
- $64 = 2^6$  byte pages.

## Virtual addresses



## Physical addresses




## TLB (4 sets, 4-way set associative)

Set	Tag	PPN	Valid	Tag	PPN	Valid	Tag	PPN	Valid	Tag	PPN	Valid
0	03	00	0	09	0D	1	00	0A	0	07	02	1
1	03	2D	1	02	10	0	04	0B	0	0A	03	0
2	02	3D	0	08	0F	0	06	05	0	03	2D	0
3	07	12	0	03	0D	1	0A	34	1	02	03	0

## Page table

VPN	PPN	Valid	VPN	PPN	Valid	VPN	PPN	Valid	VPN	PPN	Valid
00	28	1	04	10	0	08	13	1	0C	13	0
01	27	0	05	16	1	09	17	1	0D	2D	1
02	33	1	06	28	0	0A	09	1	0E	11	1
03	02	1	07	28	0	0B	01	0	0F	0D	0

# Cache (16 lines, direct mapped, physically addressed)



## Contents

Idx	Tag	Valid	Block			
0	19	1	99	11	23	11
1	15	0	01	02	11	99
2	18	1	00	02	04	08
3	36	0	11	23	11	99
4	32	1	43	6D	8f	09
5	0D	1	36	72	F0	1D
6	31	0	72	1D	F0	36
7	16	1	11	C2	DF	03

Idx	Tag	Valid	Block			
8	24	1	3A	00	51	89
9	2D	0	89	00	3A	51
A	2D	1	93	15	DA	3B
B	08	0	89	51	15	3A
C	12	0	3A	00	15	69
D	16	1	04	96	34	15
E	13	1	83	77	1B	D3
F	14	0	14	20	13	37

# Translating virtual addresses and cache lookups

translate into binary:

00 0011 1101 0100

▪ **Virtual address:** 0x03D4

▶ VPN  
TLBI  
TLBT  
TLB Hit?  
Page fault?  
PPN

▪ **Physical address:**

▶ CO  
CI  
CT  
Cache hit?  
Byte

▪ **Virtual address:** 0x0020

▶ VPN  
TLBI  
TLBT  
TLB Hit?  
Page fault?  
PPN

▪ **Physical address:**

▶ CO  
CI  
CT  
Cache hit?  
Byte

Simple memory system example

The Linux virtual memory system

Memory mapping

Dynamic allocation

# Structure of a page table

VPN	PPN	Valid	Flags
0			
1			
2			
3			
4			
	⋮		

- Showing a flat page table for simplicity.
- Actually a multi-level page table.

- Page table is accessed by MMU (hardware) so its structure is **fixed** and kept **simple**.
- Either a page is *there* or *it's not*.
  - ▶ No “page is on disk” information.
  - ▶ No “demand-paging” information.
  - ▶ If we access a non-valid page: **page fault!**
- Page faults are handled by **software** (kernel code), meaning we have **flexibility**.
  - ▶ Page fault *handler* can update the page table based on kernel data and policy.
- **But what does that actually look like?**




# VM areas

Linux organises a virtual memory space as a set of *areas*.


```
struct vm_area_struct {  
    unsigned long vm_start;  
    unsigned long vm_end;  
    pgprot_t      vm_page_prot;  
    ...  
};
```

- Each area describes the properties of a span of virtual memory.
  - ▶ May cover any number of pages.
  - ▶ Area has uniform protection/access bits (read, write, exe).
- **CPU/MMU has no idea what this is**—purely a software data structure.
  - ▶ Easier to change kernel code than hardware.

# Handling page faults




# Handling page faults




- **Segmentation fault**–page does not exist.

# Handling page faults



- **Segmentation fault**—page does not exist.
- **Normal page fault**
  - ▶ Find available physical page in RAM.
  - ▶ Fetch page contents from disk.
  - ▶ Update page table.

# Handling page faults



- **Segmentation fault**–page does not exist.
- **Normal page fault**
  - ▶ Find available physical page in RAM.
  - ▶ Fetch page contents from disk.
  - ▶ Update page table.
- **Segmentation fault**–permissions are wrong for type of access.

Simple memory system example

The Linux virtual memory system

Memory mapping

Dynamic allocation

# Memory mapping

- **VM areas are associated with disk objects**
  - ▶ Known as *memory mapping*.
  - ▶ As a programming abstraction, lets us access a file as it were memory.
  - ▶ But also a strong organising principles.
- **Area can be *backed by* (get its initial values from):**
  - ▶ **Regular file** on disk (e.g. a program executable file).
    - ▶ Initial page contents from from section of file.
  - ▶ **Anonymous file** (i.e. nothing).
    - ▶ Initial page contents are zero.
    - ▶ Once written to, like any other page.

# Memory mapping

- **VM areas are associated with disk objects**
  - ▶ Known as *memory mapping*.
  - ▶ As a programming abstraction, lets us access a file as it were memory.
  - ▶ But also a strong organising principles.
- **Area can be *backed by* (get its initial values from):**
  - ▶ **Regular file** on disk (e.g. a program executable file).
    - ▶ Initial page contents from from section of file.
  - ▶ **Anonymous file** (i.e. nothing).
    - ▶ Initial page contents are zero.
    - ▶ Once written to, like any other page.

```
struct vm_area_struct {  
    ...  
    unsigned long    vm_pgoff;  
    struct file      *vm_file;  
    ...  
};
```



# Memory mapping from userspace


```
void *mmap(void *addr, size_t length, int prot, int flags,  
           int fd, off_t offset);
```

- Map `len` bytes starting at offset `offset` of the file descriptor `fd`, preferably at address `start`.
  - ▶ `start` may be 0 for “I don't care”.
  - ▶ `prot`: `PROT_READ`, `PROT_WRITE`, `PROT_EXEC`, `PROT_NONE`.
  - ▶ `flags`: `MAP_ANONYMOUS`, `MAP_PRIVATE`, `MAP_SHARED`, many more...
- Returns a pointer to start of mapped area (may not be `start`).

```
int munmap(void *addr, size_t length);
```

- Unmaps pages starting at this address.
  - ▶ VM area is cloven in twain if we unmap pages in the middle.

```
void *mmap(void *addr, size_t length, int prot, int flags,  
           int fd, off_t offset);
```



## mmap examples

See lecture code.

Simple memory system example

The Linux virtual memory system

Memory mapping

Dynamic allocation

# User-space allocators


- `mmap()` is **powerful** but **inflexible**:
  - ▶ Smallest granularity of allocation is a page.
  - ▶ Is a system call, so fairly slow.

# User-space allocators

- `mmap()` is **powerful** but **inflexible**:
  - ▶ Smallest granularity of allocation is a page.
  - ▶ Is a system call, so fairly slow.
- Instead programmers use **dynamic memory allocators** (e.g. `malloc()`) to acquire memory at runtime.
  - ▶ Run entirely in user space—not part of the kernel.
  - ▶ Acquires memory via `mmap()` and `sbrk()`.

# User-space allocators

- `mmap()` is **powerful** but **inflexible**:
  - ▶ Smallest granularity of allocation is a page.
  - ▶ Is a system call, so fairly slow.
- Instead programmers use **dynamic memory allocators** (e.g. `malloc()`) to acquire memory at runtime.
  - ▶ Run entirely in user space—not part of the kernel.
  - ▶ Acquires memory via `mmap()` and `sbrk()`.
- Region of virtual memory managed by such an allocator is known as *the heap*.
  - ▶ No relation to the datastructure known as a heap.
  - ▶ May have multiple heaps; heap might not be contiguous.
  - ▶ When we say *the heap*, we mean whatever `malloc()` manages by default.



# Dynamic memory allocation

- Allocator maintains heap as collection of *blocks* of varying size, each of which is either *allocated* or *free*.
- Types of allocators:
  - Explicit allocator:** application manually allocates and frees space.
    - ▶ E.g. `malloc()` in C.
  - Implicit allocator:** application allocates but freeing is automatic.
    - ▶ E.g. garbage collection in F#, SML, Haskell, Futhark, and Lisp.
- CompSys discusses simple explicit allocators.
  - ▶ **Implicit allocators** may be touched upon in PLD.



# Allocator API

```
void *malloc(size_t size);
```

- Returns a pointer to a memory block of *at least* size bytes.
  - ▶ May round up to ensure address is always aligned.
- NULL on failure.

```
void free(void *ptr);
```

- Returns block of memory to heap. return to the bottom of heap, which means free the heap
- p must have been returned by malloc() (or one of the variants).

```
void *realloc(void *ptr, size_t size);
```

```
void *calloc(size_t nmemb, size_t size);
```

- Less crucial, but sometimes useful.

# System-level building blocks

**malloc()** uses the following functions to request memory from the kernel.

```
void *mmap(void *addr, size_t length, int prot, int flags,  
           int fd, off_t offset);
```

- Use `MAP_ANONYMOUS` to map fresh pages (will be demand-paged when first accessed).


```
void *sbrk(intptr_t increment);
```

- Modifies *heap pointer* (traditionally called the *program break*).
- Maps or unmaps new anonymous read-write pages as necessary.
- Similar to growing the stack by decrementing stack pointer.
  - ▶ But with a *system call*: there is no register for the heap pointer.

This memory is then split into blocks and returned to callers of **malloc()** as needed.

# Notation used in the following slides

- We show how the memory managed by `malloc` is split into blocks.
- We show 4-byte words instead of single bytes.



Free word

Allocated word (will use different colours to indicate different allocated blocks)

# Allocation example

```
p1 = malloc(4*4);
```




# Allocation example

```
p1 = malloc(4*4);
```



```
p2 = malloc(5*4);
```




# Allocation example


```
p1 = malloc(4*4);
```



```
p2 = malloc(5*4);
```



```
p3 = malloc(6*4);
```




# Allocation example


```
p1 = malloc(4*4);
```




```
p2 = malloc(5*4);
```



```
p3 = malloc(6*4);
```



```
free(p2);
```




# Allocation example


```
p1 = malloc(4*4);
```




```
p2 = malloc(5*4);
```



```
p3 = malloc(6*4);
```



```
free(p2);
```



```
p4 = malloc(2*4);
```





# Constraints

## Applications

- Can issue arbitrary sequence of `malloc()` and `free()` requests.
- `free()` must be to a `malloc()`ed block.

## Allocators

- Can't control number or size of allocated blocks.
- Blocks must be *contiguous* in memory.
- Must respond immediately to `malloc()` requests.
  - ▶ i.e. cannot reorder or delay requests.
- Must allocate blocks from free memory.
  - ▶ i.e. blocks cannot overlap.
- Can manipulate and modify only free memory.
  - ▶ When a block has been returned from `malloc()`, *that memory belongs to the application*.
- Cannot move blocks once `malloc()`d.
  - ▶ ...can someone say why?

# Performance Goal I: *throughput*

- We are given some sequence of **malloc** and **free** requests:
  - ▶  $R_0, R_1, \dots, R_{n-1}$ .
- Goal: maximise throughput
- Throughput:
  - ▶ Number of completed requests per unit time.
    - ▶ E.g. suppose 5000 `malloc()`s and 5000 `free()`s in 10 seconds.
    - ▶ Throughput is 1000 ops/s.

## Performance Goal II: *peak memory utilisation*

- We are given some sequence of **malloc** and **free** requests:

▶  $R_0, R_1, \dots, R_{n-1}$ .

- Goal: maximise memory utilisation

Aggregate payload  $P_k$ :

- `malloc(p)` results in a block with a *payload* of  $p$  bytes.
- $P_k$  is the sum of currently allocated payloads after request  $R_k$ .

Current heap size  $H_k$ :

- Total amount of heap memory requested from the system.
- Grows only when using `sbrk()` or `mmap()`.

Peak memory utilisation after  $k$  requests  $U_k$ :

- $U_k = \max_{i < k} \frac{P_i}{H_k}$
- Indicates how much memory we have *reserved* without *using*.


defragmentation: press  
all the discrete data  
block together to get  
more continuous memory

# Fragmentation

- **Poor memory utilisation is caused by**
  - ▶ *internal* fragmentation, and
  - ▶ *external* fragmentation.

# Internal fragmentation

*Internal fragmentation* occurs when the payload is smaller than the block size.




- Causes**
- Overhead of maintaining heap data structures.
  - Padding for alignment purposes.
  - Explicit policy decisions, e.g. returning a bigger block than necessary because it's faster.

Depends only on *previous* requests, so easy to measure.


# External fragmentation

Occurs when there is enough aggregate heap memory, but no single free block is large enough.

```
p1 = malloc(4*4);
```




```
p2 = malloc(5*4);
```



a rule of thumb:  
always allocate  
bigger block

```
p3 = malloc(6*4);
```



```
free(p2);
```



one direction,  
不能从尾巴接到头

```
p4 = malloc(6*4);
```

**No free block with room for six words.**


Depends on future requests, so difficult to measure.

# Key implementation questions

- How does `free()` know the size of the block given just a pointer?
- How do we keep track of free blocks?

# Knowing how much to free (standard method)

- Keep the length of a block in the word preceding the block.
  - ▶ This word is often called the *header field* or *header*.
  - ▶ We can also use more words to store even more metadata if we wish.
- Requires an extra word for every allocated block.



`free(p0)`





# Keeping track of free blocks

- Method 1: *Implicit list* using length—links all blocks.



- Method 2: *Explicit list* among the free blocks, using pointers or offsets.



- Method 3: *Segregated free list*, different free lists for different sizes of blocks.

# Conclusions

- `mmap()` syscall allows processes to map virtual memory.
  - ▶ Can map files to memory, or make anonymous mapping.
  - ▶ Can share memory between processes.
- `malloc()` is a *userspace* memory manager.
  - ▶ Not a system call itself.
  - ▶ Requests memory from kernel with `mmap()` and `sbrk()` and then parcels it out.
  - ▶ *Internal fragmentation* is when allocated blocks have wasted space.
  - ▶ *External fragmentation* is when free space is split into many small blocks.